

Saneesh Morais

London, UK — Open to UK hybrid roles — +44 7466 616145
saneeshmorais@gmail.com — LinkedIn: Saneesh Morais — Portfolio — GitHub

Personal Profile

First Class Honours graduate in BSc (Hons) Computing with Artificial Intelligence Technology, focused on building practical AI features, LLM integrations, agents, prompt flows, APIs, automation tools, and safe AI workflows. Built working AI applications using Python, JavaScript, TypeScript, FastAPI, REST APIs, SQL, SQLite, LLM APIs, RAG, embeddings, structured outputs, Git, Docker, and GitHub Actions. Strong at quickly prototyping ideas, testing model outputs, documenting reusable patterns, and turning business problems into AI-powered features that are useful, safe, and maintainable.

Technical Skills

AI Engineering: OpenAI API, Anthropic API, Claude, LLM Agents, Prompt Engineering, Prompt Flows, Structured Outputs, Tool Use

Search and AI Features: RAG, Embeddings, Smart Search, Summarisation, Document Q&A, Recommendations, Chatbot Workflows

Programming and APIs: Python, JavaScript, TypeScript, FastAPI, Flask, REST APIs, JSON, SQL, HTML, CSS

Data and Evaluation: pandas, NumPy, SQLite, PostgreSQL, ChromaDB, Data Validation, Test Cases, Spot Checks, Feedback Loops

Development Workflow: Git, GitHub, Docker, GitHub Actions, CI/CD, Debugging, Refactoring, Documentation, API Testing

Governance and Safety: Output Validation, Human Review, Risk Scoring, Guardrails, Audit Logs, Failure Analysis, AI Limitations

Education

BSc (Hons) Computing with Artificial Intelligence Technology **Oct 2023 – Jun 2026**
Northumbria University London First Class Honours

Relevant Study: Programming for AI, Software Development, Machine Learning, Data Science, Cloud Computing, AI Ethics and Governance.

Work Experience

Data Analyst Intern **Jan 2025 – Mar 2025**
Savant Technologies, London – Remote

- Developed Python and SQL workflows across 10K+ operational records, covering data cleaning, validation, transformation, analysis, and reporting.
- Built reusable processing scripts and dashboard outputs, reducing manual reporting time by 15%.
- Added validation checks for missing, duplicate, and inconsistent records, reducing reporting inconsistencies by 10%.
- Documented assumptions, data issues, workflows, and improvement recommendations for stakeholder review.

AI Intern **May 2024 – Aug 2024**
3NS.ai, London – Hybrid

- Built and tested AI prototypes for financial analysis, improving forecasting accuracy from 72% to 85%.
- Evaluated 20+ variables through feature engineering, structured testing, output validation, and error analysis.
- Developed an NLP document-review workflow using TF-IDF and cosine similarity, reducing manual review time by 30%.
- Recorded model behaviour, limitations, outputs, assumptions, and improvements to support reliable AI adoption.

Selected Projects

AgentGuard OS – Governed AI Agent Platform [Video](#)

- Built a FastAPI and SQLite platform that evaluates AI agent requests, detects unsafe actions, scores risk, and blocks unreliable outputs.
- Designed prompt checks, structured JSON decisions, approval gates, audit logs, cost guardrails, circuit breakers, and failure-handling controls.
- Tested legitimate and adversarial requests to compare outputs, identify edge cases, and improve safety, reliability, and escalation behaviour.

Iris Job Assistant – Full-Stack AI Workflow Application [Link](#)

- Built a FastAPI, SQLite, HTML, CSS, and JavaScript application integrating LLM APIs into multi-step, human-reviewed workflows.
- Developed job analysis, summarisation-style outputs, structured application packets, ATS checks, document generation, and quality gates.
- Used prompt design, structured outputs, test cases, validation checks, and failure handling to prevent weak AI outputs progressing automatically.

RAG Pipeline from Scratch [Link](#)

- Built a Python RAG pipeline covering document loading, chunking, embedding generation, ChromaDB storage, retrieval, and grounded answer generation.
- Tested retrieval quality using different chunk sizes, similarity thresholds, prompt formats, and query patterns to improve answer relevance.
- Documented prompt flows, weak matches, retrieval failures, model limitations, and improvements for more reliable smart-search features.

Certifications

IBM Artificial Intelligence Fundamentals (2025) — Databricks Generative AI Fundamentals (2025) — Python Programming – Networkz Systems (2023)

Languages and Interests

Languages: English — Tamil — Hindi — Malayalam — Spanish – Beginner — French – Beginner

Interests: LLM applications, RegTech, smart search, AI agents, automation, safe AI outputs, product engineering, chess.